

# LLM EVALUATION AND SAMPLE EFFICIENT PREFERENCE MEASUREMENT

Minha Hwang



# AGENDA

LLM 101: 3 TRAINING STAGES

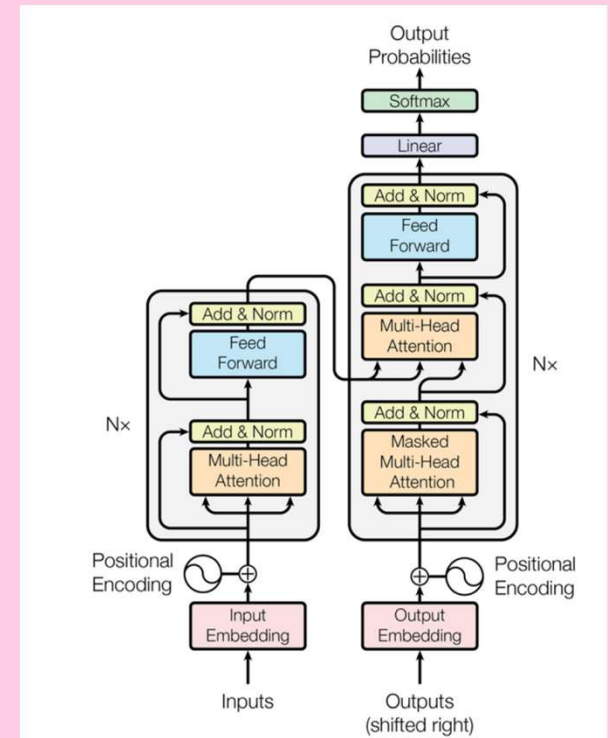
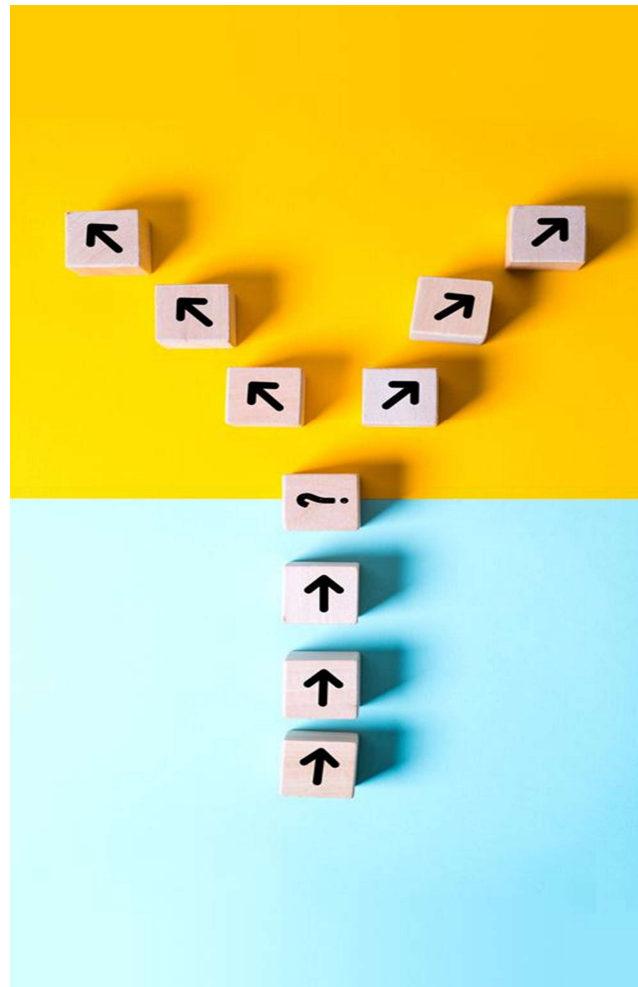
PRE-TRAINING EVAL

POST-TRAINING EVAL

RAG EVAL

MEASUREMENT THEORY

MAX DIFF





## **LLM 101:3 TRAINING STAGES**

## AUTOREGRESSIVE LANGUAGE MODELS: PROBABILITY KERNEL

- A language model is a **probability kernel  $\mu$**  given a prefix of words:  **$\mu: X \rightarrow Pr(Y)$** 
  - Stochastic in nature: **A same prefix  $X$**  can give a **random output** sampled from a probability distribution  $\mu_X$  (i.e., generative) → A key reason for factual inaccuracy, inconsistency or hallucination (making stuff up)
- A language model calculates  **$Pr(s)$**  given a sequence of words:  **$s = (w_1, w_2, \dots, w_{T-1}, w_T)$**
- An autoregressive language model calculates this **conditional on a previous sequence of words**:

$$\begin{aligned} Pr(s) &= Pr(w_1, w_2, \dots, w_{T-1}, w_T) \\ &= \prod_{t=1}^T Pr(w_t | w_1, w_2, \dots, w_{t-1}) \end{aligned}$$

- **Next-word prediction**: Given a prefix  $(w_1, w_2, \dots, w_{t-1})$ , calculate the probability of the next word  $w_t$  (Conceptually same to time series with path dependence)

Source: Prof. Kyunghyun Cho

## AUTOREGRESSIVE LANGUAGE MODELS: SIMPLE EXAMPLE

- 4-word sentence example: “I am a student”

$$Pr(s) = Pr(w_1, w_2, w_3, w_4) = Pr(w_1) \times Pr(w_2|w_1) \times Pr(w_3|w_1, w_2) \times Pr(w_4|w_1, w_2, w_3)$$

- All you need is “**counting**” (if there are large amounts of data)

$$Pr(w_2|w_1) = \frac{count(w_1, w_2)}{count(w_1)} \longrightarrow \text{2-grams (Bigrams)}$$

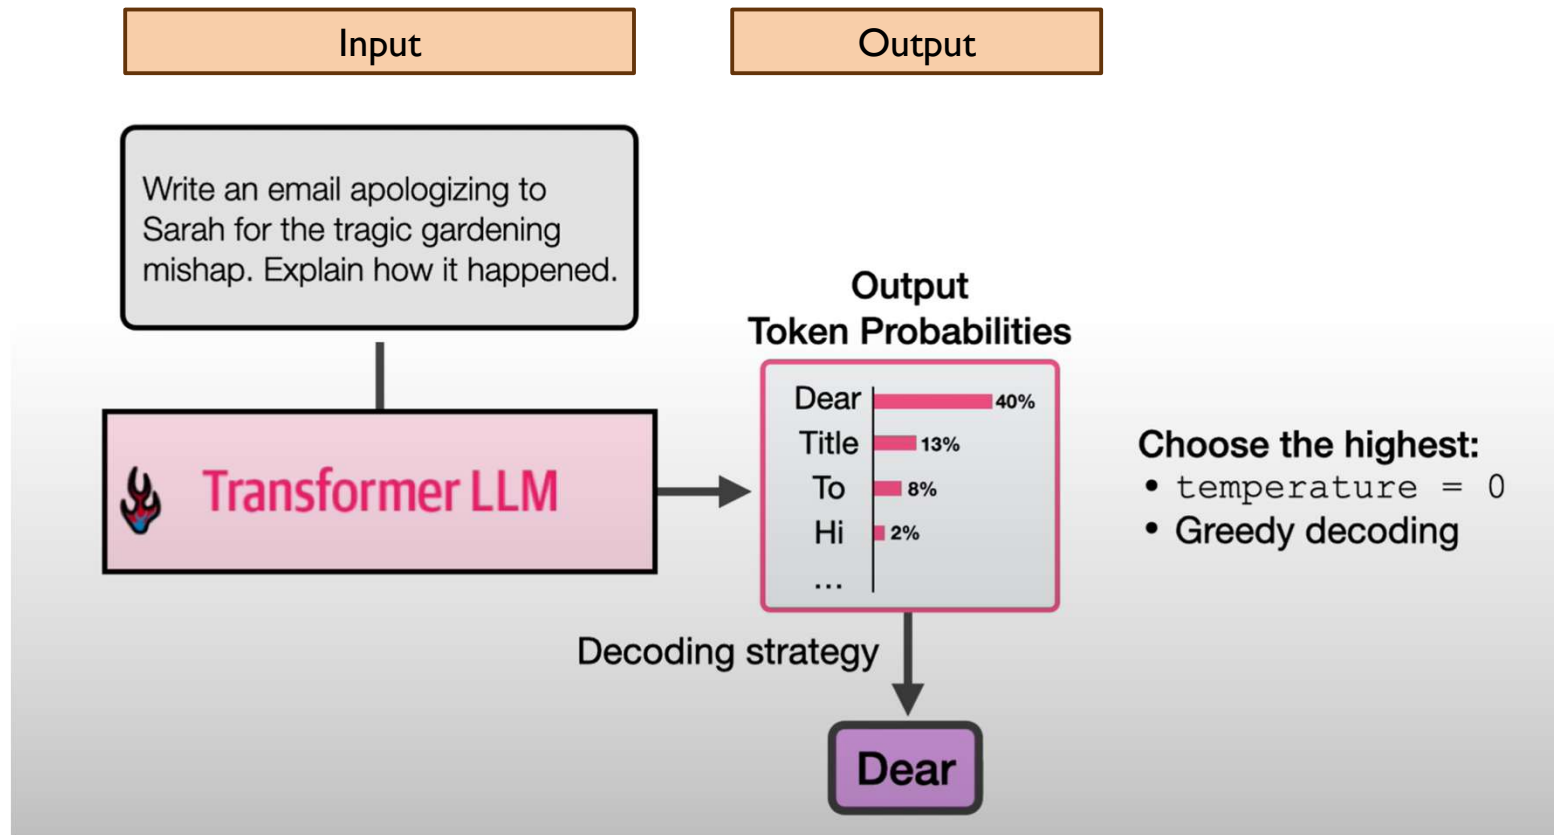
$$Pr(w_3|w_1, w_2) = \frac{count(w_1, w_2, w_3)}{count(w_1, w_2)} \longrightarrow \text{3-grams (Trigrams)}$$

$$Pr(w_4|w_1, w_2, w_3) = \frac{count(w_1, w_2, w_3, w_4)}{count(w_1, w_2, w_3)} \longrightarrow \text{4-grams}$$

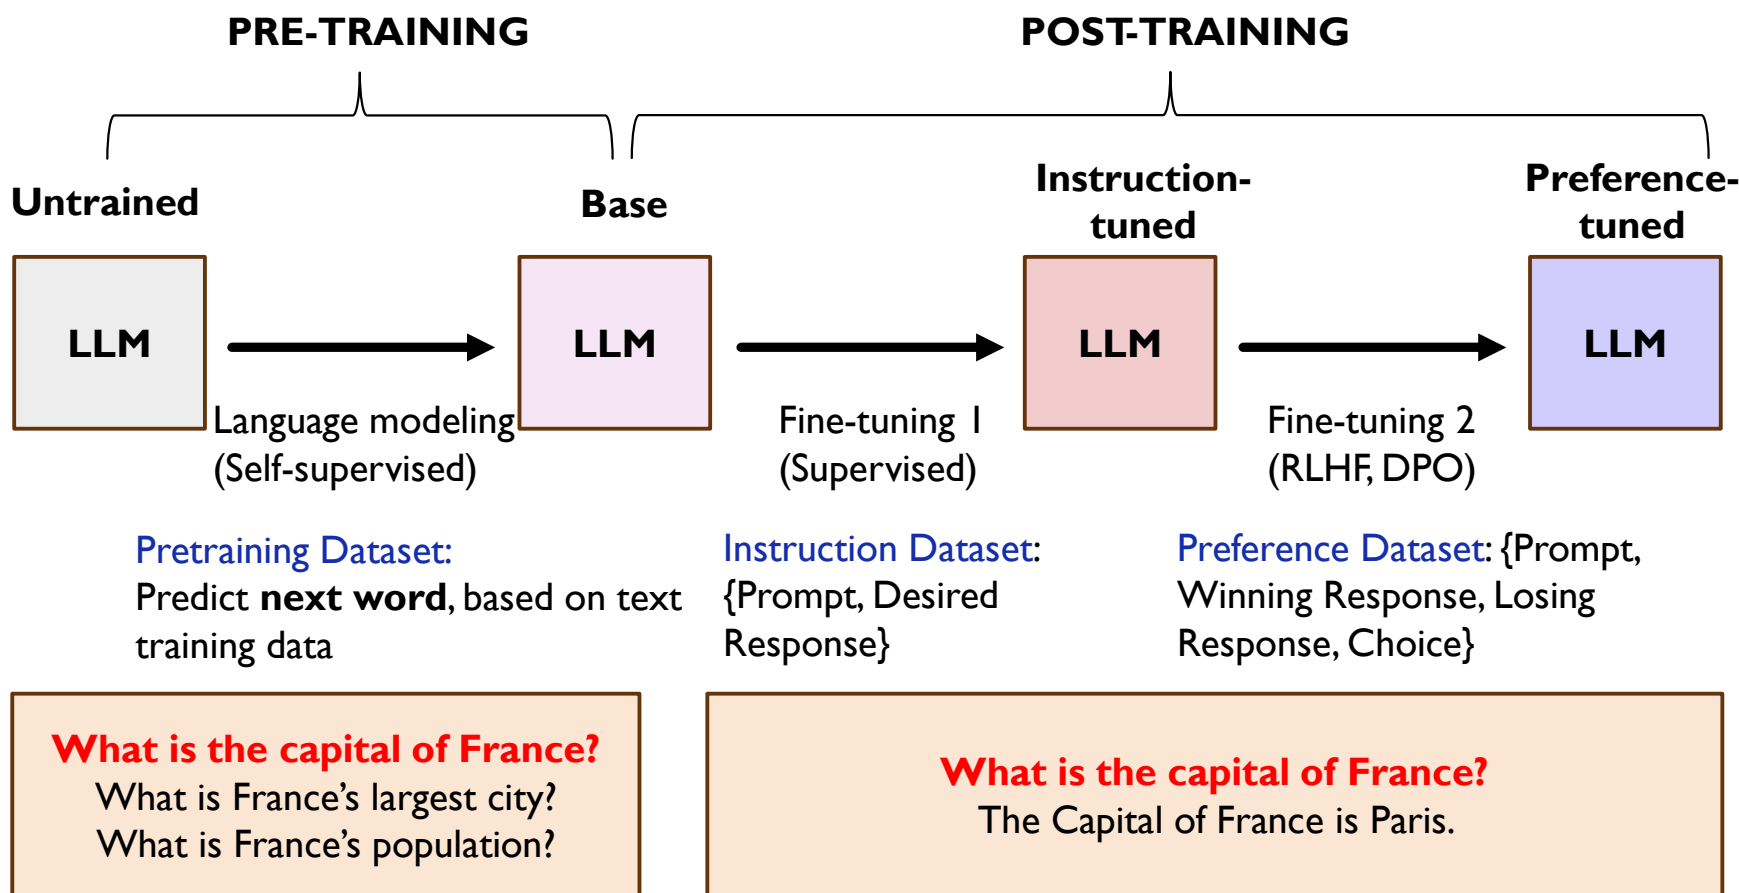
- Problems:
  - This requires **a lot of space (RAM)**
  - Count-based language models **cannot generalize**: A certain sentence **does not appear** in the corpus

Source: Prof. Kyunghyun Cho

# TRANSFORMER LLM: INPUT AND OUTPUT



# LLM TRAINING: 3 STEP RECIPE



Source: Minha Hwang; [ChatGPT Prompt Engineering for Developers - DeepLearning.AI](#)

# REASONING LLMS: DeepSeek-R1 EXAMPLE

**Prompt:** Solve this logic puzzle: If  $A > B$  and  $B > C$ , who is greatest?

**Response:** Thought for 9 seconds ▾

🔍 First, the puzzle states: If  $A > B$  and  $B > C$ , who is greatest?

I need to understand what this means. It seems like a simple logic problem involving comparisons. The symbols " $>$ " probably mean "is greater than." So, A is greater than B, and B is greater than C.

In terms of greatness, if A is greater than B, and B is greater than C, then A should be greater than C as well. That means A is the greatest among A, B, and C.

This is a straightforward transitive property. If  $A > B$  and  $B > C$ , then  $A > C$ . So, A is the greatest.

The question is "who is greatest?" but it doesn't specify what A, B, and C are. In logic puzzles, sometimes A, B, C could represent people or things. But here, since it's a logic puzzle, I think we're to assume that A, B, and C are entities being compared based on their values.

Therefore, based on the given inequalities, A is the greatest.

I think that's the answer.

Based on the given conditions:

- $A > B$  (A is greater than B)
- $B > C$  (B is greater than C)

By the transitive property,  $A > C$ . Therefore, A is the greatest among A, B, and C.

A conversation between User and Assistant. The user asks a question, and the Assistant solves it. The assistant first thinks about the reasoning process in the mind and then provides the user with the answer. The reasoning process and answer are enclosed within `<think>` `</think>` and `<answer>` `</answer>` tags, respectively, i.e., `<think>` reasoning process here `</think>` `<answer>` answer here `</answer>`. User: **prompt**. Assistant:

Table 1 | Template for DeepSeek-R1-Zero. **prompt** will be replaced with the specific reasoning question during training.



---

## REASONING MODEL (TEST-TIME COMPUTE)

Normal LLM:



Reasoning LLM:



<https://newsletter.languagemodels.co/p/the-illustrated-deepseek-r1> (together.ai)



## **PRE-TRAINING EVAL**

## PRE-TRAINED EVAL - PERPLEXITY: PREDICTIVE ACCURACY (1/2)

**Pre-Training:** How well a language model **predicts a sequence of words**.

- **Low perplexity: better predictive performance, more capabilities**
- **Not a good measure** to evaluate model that have been **post-trained**
- **Perplexity of 3** – Interpretation: This model has a **1 in 3 chance of predicting next token correctly**.
- More structured data: more predictable, lower perplexity
- Bigger vocabulary: higher perplexity
- Longer context window: lower perplexity
- Exponentiation of the average negative log-likelihood of the predictive probability of each word in a test dataset

$$\text{Perplexity(PP)} = \exp\left(-\frac{1}{N} \sum_{i=1}^N \log P(w_i | w_{<i})\right)$$

The students open their

Language Model

Input

Books

Laptops

Vocab

Output

## PRE-TRAINED EVAL - PERPLEXITY: PREDICTIVE ACCURACY (2/2)

**Example Calculation:** “The cat sat on the mat”

- $P(\text{“The”}) = 0.2$
- $P(\text{“cat”} | \text{“The”}) = 0.1$
- $P(\text{“sat”} | \text{“The cat”}) = 0.15$
- $P(\text{“on”} | \text{“The cat sat”}) = 0.3$
- $P(\text{“the”} | \text{“The cat sat on”}) = 0.25$
- $P(\text{“mat”} | \text{“The cat sat on the”}) = 0.05$

$$\text{Perplexity(PP)} = \exp\left(-\frac{1}{N} \sum_{i=1}^N \log P(w_i | w_{<i})\right)$$

First, calculate the average negative log probability:

$$-\frac{1}{6} (\log(0.2) + \log(0.1) + \log(0.15) + \log(0.3) + \log(0.25) + \log(0.05)) \simeq 1.8992$$

Then, exponentiate to find perplexity:

$$PP = \exp(1.8992) \simeq 6.68$$

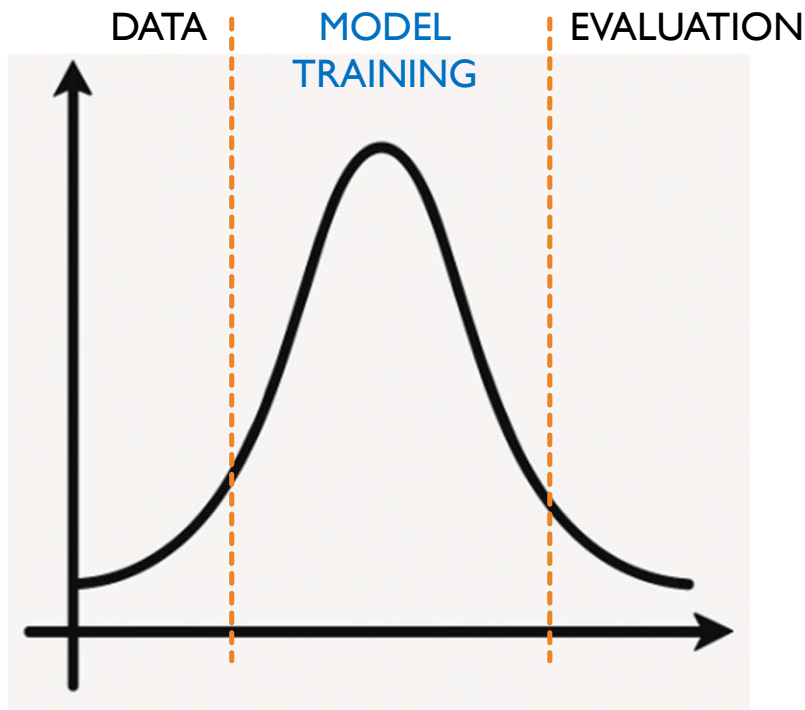
This means that the model, on average, considers about 6.68 possible next words, indicating its uncertainty in prediction.



## **POST-TRAINING EVAL**

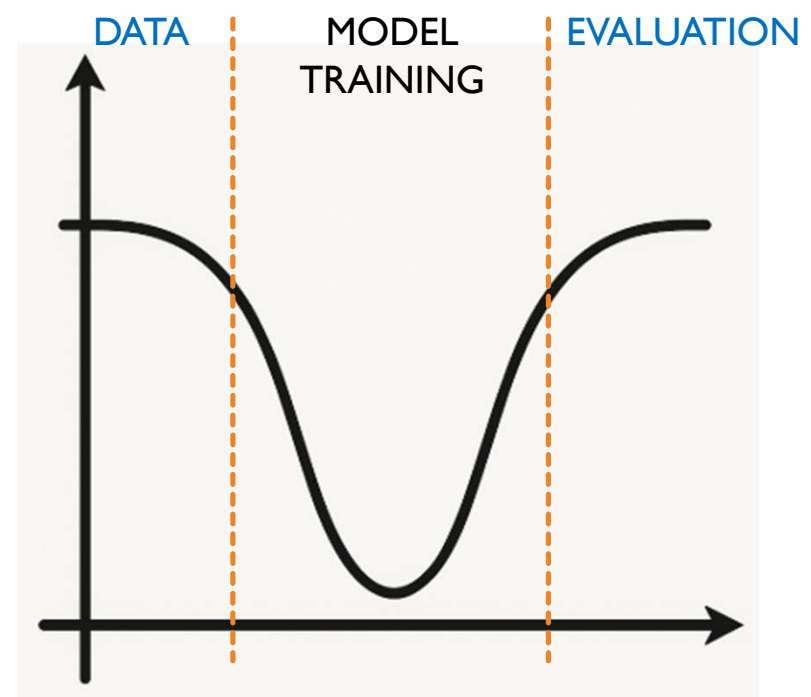
# SHIFT IN RELATIVE EFFORTS IN AI/ML ENGINEERING

## Traditional ML



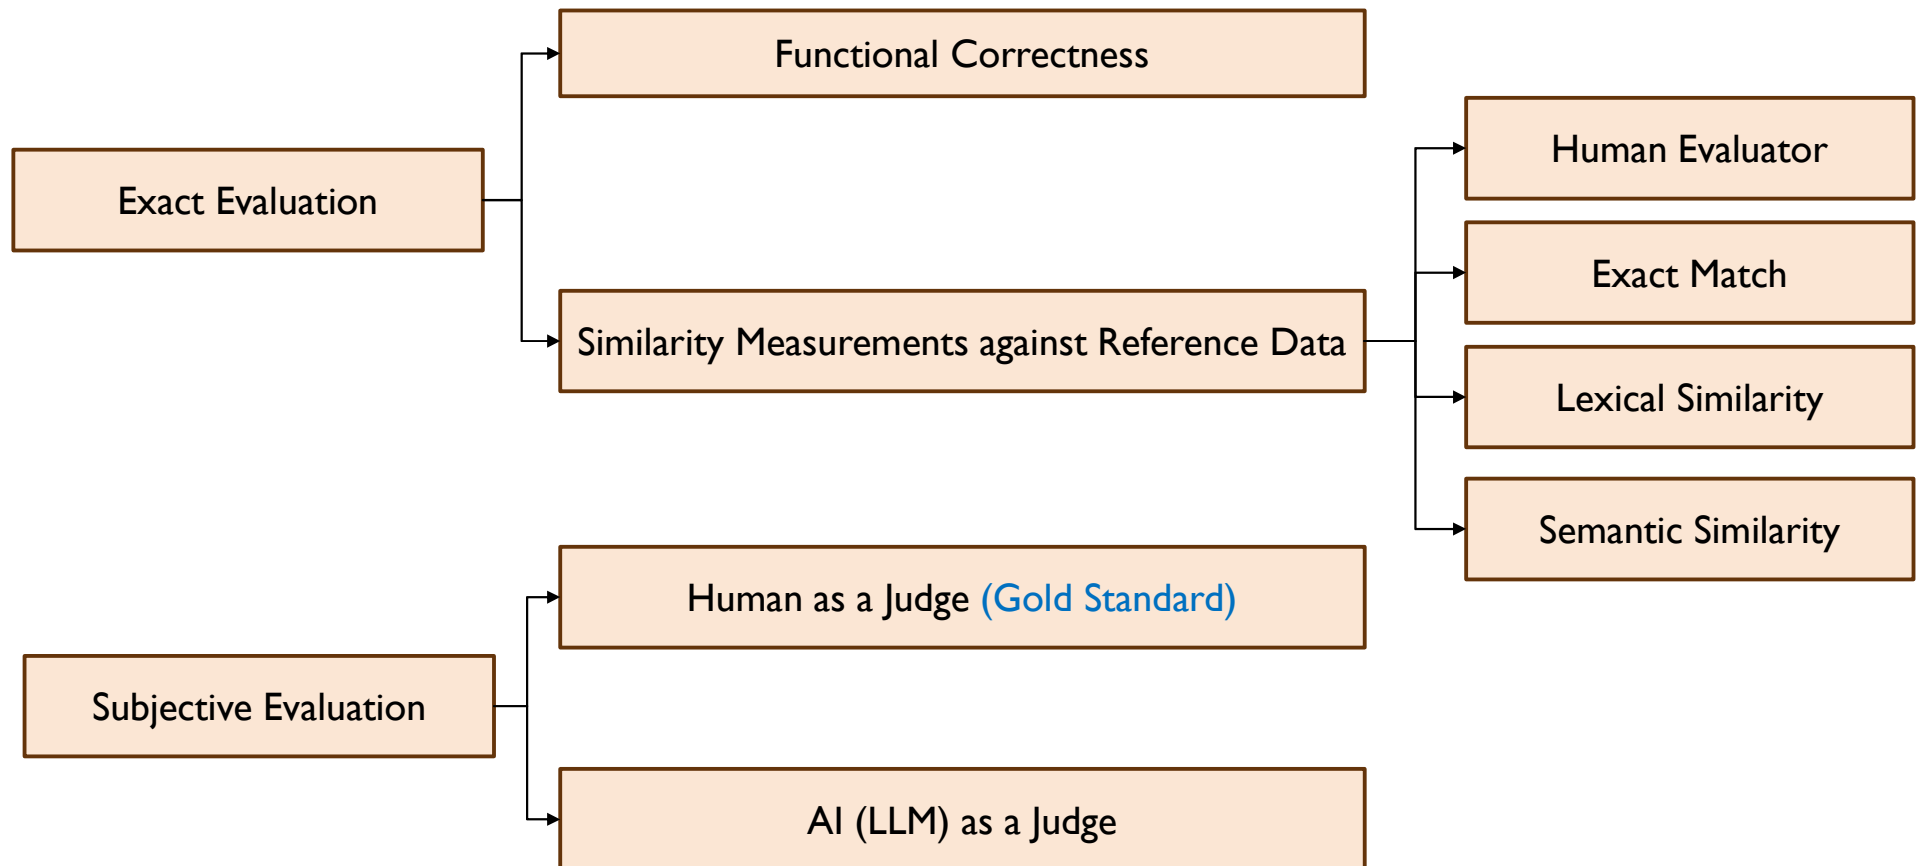
- **Concentrate on model training:** recommendation, ranking
- **Train models from scratch:** model architecture, hyper-parameter optimization and training recipes

## Generative AI / LLM



- Emphasize **high quality data curation** and **model evaluation**
- **Pre-trained model:** limited adaptation (fine-tuning)

# POST-TRAINED MODEL EVALUATION: OVERVIEWS



# POST-TRAINED MODEL EVALUATION (1/3): EXACT EVALUATION

## Functional Correctness

### What it is

- Measures whether your application does what it is intended to do
- Not always straightforward to measure
- Can be automated for certain tasks: code generation, math

### Examples

- **Code generation capabilities:** Automated functional correctness measurements (e.g., unit tests) – OpenAI's HumanEval, Google's MBPP, Spider (text-to-SQL)
- **Mathematical problem-solving, Fact-based Q&A**

### How to measure

- For each coding problem, a number of code samples (i.e.,  $k$ ) are generated
- **pass@k**: A fraction of solved problems out of all problems

## Similarity Measurements against Reference Data

### What it is

- Evaluate AI's outputs against reference data – ground truths (reference-based)
- Reference data: (input, reference responses) - multiple reference responses possible
- Generated responses that are more similar to the reference responses are considered better
- Bottlenecked by how much and how fast reference data can be generated (human or AI)

### Examples

- **BLEU, ROUGE, METEOR++, TER, BERTScore**
- Common for **translation tasks**

### How to measure

- 4 ways to measure similarity: **human evaluator**, **exact match**, **lexical similarity** (fuzzy matching, edit-distance, counting how many tokens overlap), **semantic similarity** (embedding-based)



# POST-TRAINED MODEL EVALUATION (2/3): EXACT EVALUATION

## Functional Correctness

- **Pass@k:** or code evaluation (k: number of trial)
- 10 problems and a model solves 5 with k=3: pass@3 score of 50%
- The more code samples a model generates, the more chance the model has at solving each problem, hence the greater the final score
- In expectation, pass@1 score should be lower than pass@10
- Unit tests: run automated tests against generated outputs to check correctness
- Game bots: what score it gets in playing “Tetris”
- Logical validation: Ensure responses follow logical consistency (e.g., in reasoning tasks)

## Similarity Measurements against Reference Data

### BLEU (Bilingual Evaluation Understudy)

- Measures n-gram overlap between generated and reference text up to length N (usually 4)
- Used in machine translation and text generation tasks.
- Formula

For each  $n \in \{1, \dots, N\}$ ,

$$p_n = \frac{\sum_{\text{ngram} \in \text{Cand}_n} \text{clip}(\text{Count}_{\text{Cand}}(\text{ngram}), \max_i \text{Count}_{\text{Ref}_i}(\text{ngram}))}{\sum_{\text{ngram} \in \text{Cand}_n} \text{Count}_{\text{Cand}}(\text{ngram})},$$

- BP: Brevity Penalty (short sentences)
  - $c$  = length of the candidate translation
  - $r$  = effective reference length (closest in length to  $c$ )

$$\text{BP} = \begin{cases} 1, & c > r \\ \exp\left(1 - \frac{r}{c}\right), & c \leq r \end{cases}$$
$$\text{BLUE} = \text{BP} \times \exp\left(\sum_{n=1}^N w_n \cdot \log p_n\right)$$

+ Simple, language-agnostic, fast

- Ignores meaning, synonyms/paraphrases not counted, reward exact wording

- Example: If the reference text is “The cat is on the mat” and the generated text is “The cat is sitting on the mat.,” BLEU calculates overlapping words and assigns a score.

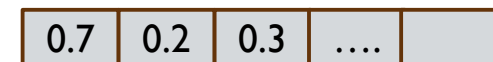
---

# SENTENCE EMBEDDING CAPTURES SEMANTIC MEANING AND CONTEXT-AWARE

**Sentence Embeddings:** Dense vector representations of sentence that capture their semantic meaning. Context-aware

**Generation** “I have a dream that one day this nation will rise up and live out the true meaning of its creed: We hold these truths to be self-evident, that all men are created equal.”

**Sentence**



**Embedding Vector**

# POST-TRAINED MODEL EVALUATION (3/3): EXACT EVALUATION

## Similarity Measurements against Reference Data

### Edit Distance (Levenstein Distance)

- Measures the number of insertions, deletions, and substitutions needed to convert one string to another.
- Character-level comparisons

- Example: “hello” → ”helo” (1 edit),  
“hello” → ”world” (4 edits)

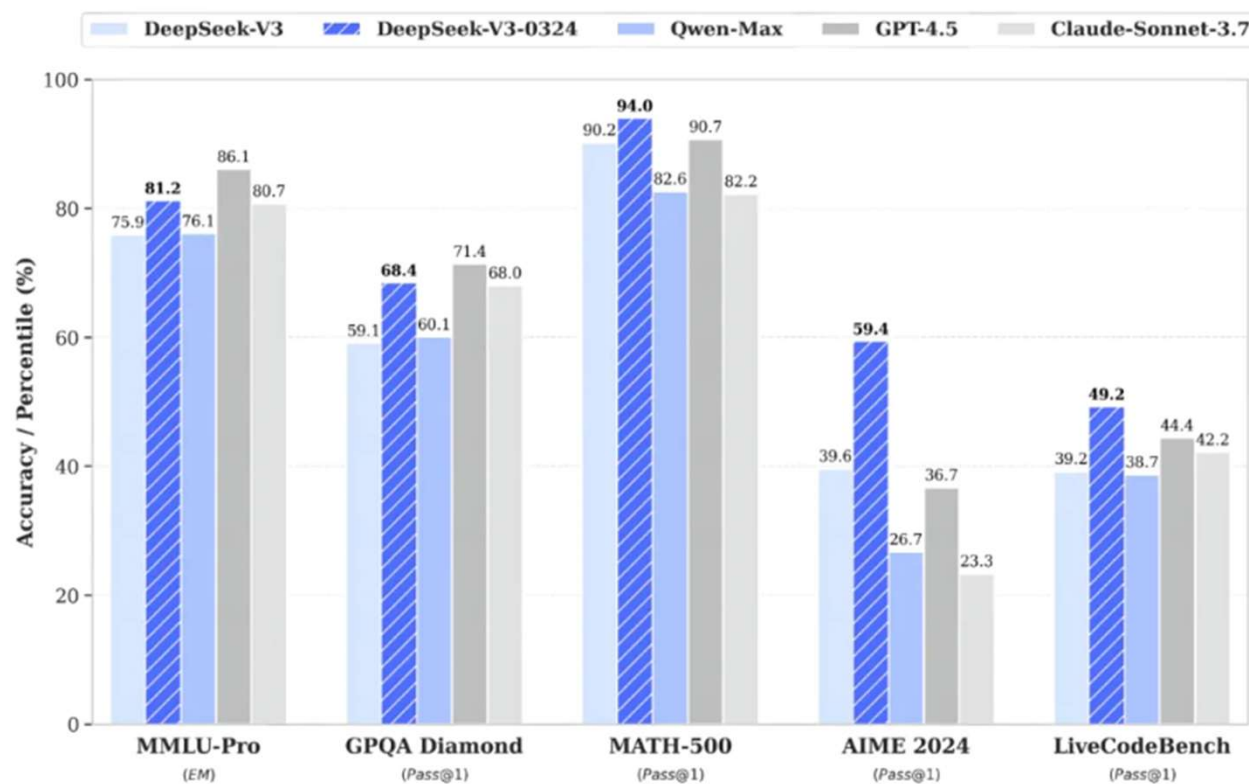
### Embedding-based Similarity (Cosine Similarity)

- Measures semantic similarity by comparing word embeddings in vector space
- Used in paraphrasing and open-ended text generation
- Formula:

$$\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$

- Example: If “The **cat** is on the **mat**” and “A **feline** is resting on a **rug**” have similar vector representations, they are considered semantically similar

# POPULAR REASONING BENCHMARKS



- **MMLU-Pro**: Massive Multitask Language Understanding – language comprehension and reasoning
- **GPQA Diamond** (Google-Proof Q&A – Hard): 198 post-graduate multiple-choice questions in physics, chemistry, and biology – expert-level
- **MATH-500**: A subset of the original 12k-question MATH benchmark (Competition-level high-school)
- **AIME 2024 (MATH)**: 30 problems from the 2024 American Invitational Mathematics Examination
- **LiveCodeBench**: > 300 coding problems continuously scraped from recent LeetCode, AtCoder, and Codeforces contests (May 2023 – Feb 2024)

# POST-TRAINED MODEL EVALUTION: SUBJECT EVALUATION (HUMAN PREFERENCE)

## Example Criteria

### Satisfaction

- **Satisfaction**: Which response makes the user satisfied with the results?
  - **Task Completion (+)**: Was the user able to complete his/her tasks?
  - **Efforts (-)**: How much efforts did the user put?
- **Helpfulness**: Which response answers the query and follow the specified instructions better?

### Factuality (Against Hallucination)

- **Factuality**: Which response provides more accurate answers without hallucinations, even for questions that require very precise or niche knowledge?

### Freshness (Against Stale Knowledge)

- **Freshness (inspired by [FreshLLMs](#))**: Which response contains more up-to-date information? A model excels in this criterion if it is able to answer queries with “fresh” information.

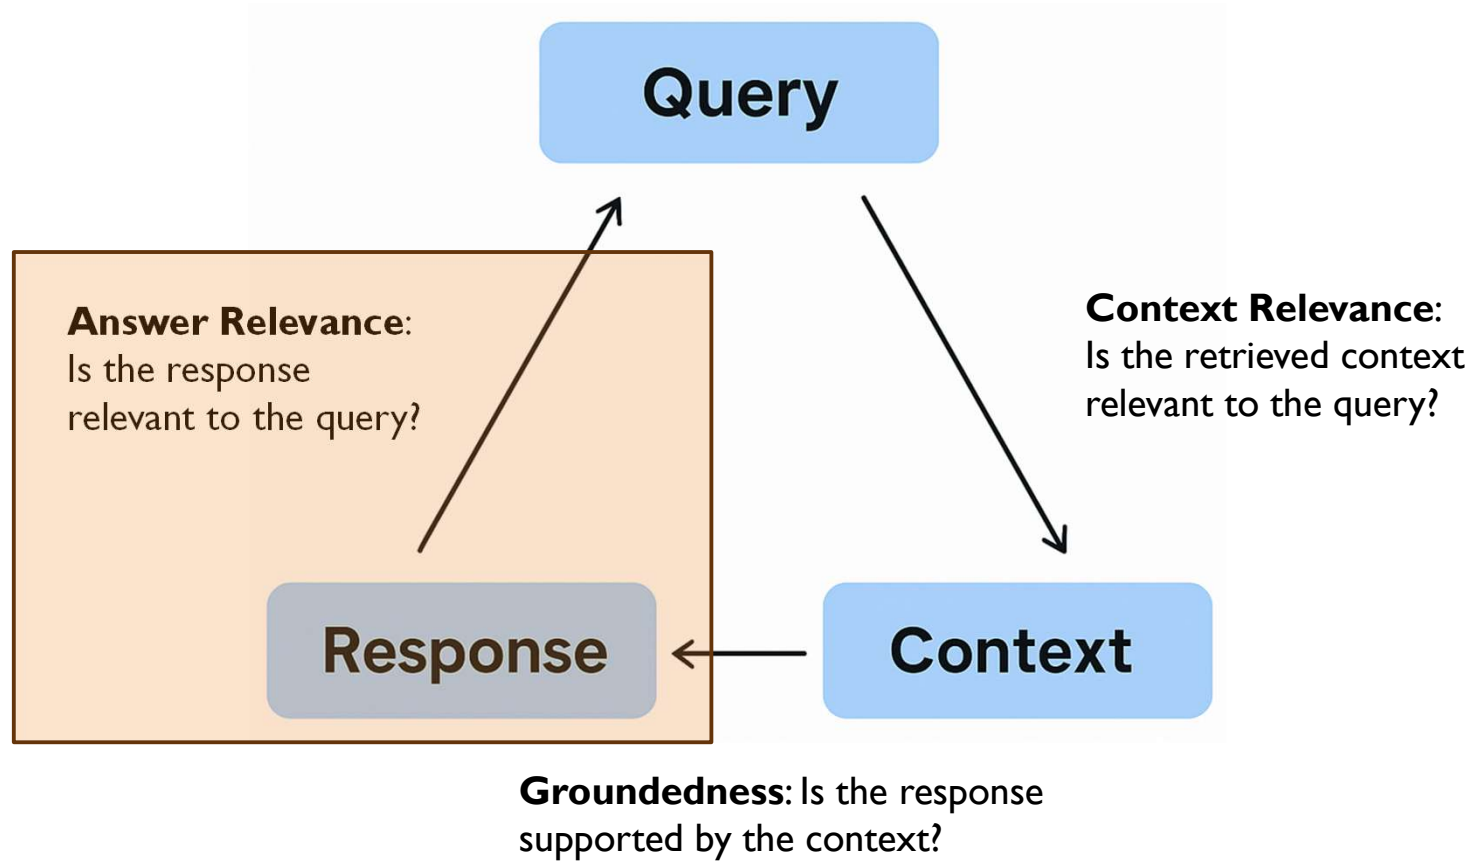
Source: [Introducing PPLX Online LLMs](#)  
[Perplexity enhances LLMs with holistic quality evaluation](#)

**Other Criteria**: Relevance, Coherence, Fluency, Conciseness, Harmfulness, Maliciousness, Controversiality, Misogyny, Insensitivity, Criminality



## **RAG EVAL**

## THE RAG TRIAD: EVALUATING RAG SYSTEM



Source: (1) [Evaluating RAG Systems: A Comprehensive Approach to Assessing Retrieval and Generation Performance](#) | LinkedIn

## EVALUATING RAG SYSTEM: ANSWER RELEVANCE (1/2) – QUERY/ANSWER

Generation component: synthesizes responses based on the retrieved data

### Answer Relevance

- Evaluates the **relevance of the generated answer** to the user's (original) query, **penalizing answers** that are **incomplete** or contain **redundant** information.
    - How relevant is the generated answer to the question?
    - The metric does not consider factual accuracy but rather focuses on the **directness and appropriateness** of the response in addressing the question.
- 
- Evaluation process: **Similarity-based** - Quantify the relevance using BERT-style models, sentence embedding models, or LLMs. The metric is computed in a reference-free manner, using the mean cosine similarity between multiple questions generated from the answer and the original question.
    - 1. Prompt the LLM to generate multiple appropriate questions based on the generated answer.
    - 2. Calculate the cosine similarity between the embedding of each generated question and the original question
    - 3. Compute the mean of these cosine similarity scores to obtain the final answer relevance score.

*Answer Relevance*

$$= \frac{\sum \text{cosine\_similarity}(\text{embedding of original question}, \text{embedding of question generated from the answer})}{\text{Total number of generated questions}}$$



## EVALUATING RAG SYSTEM: ANSWER RELEVANCE (2/2) – QUERY/ANSWER

### Example

- Question: **“Where is France and what is its capital?”**
- Answer:
  - **High relevance** “France is in western Europe, and Paris is its capital.” (**directly and completely address** both parts of the original question.)
  - **Low relevance**: “France is in western Europe.” (only **partially addresses** the original query)

Answer relevance metric ensures that the answers provided by RAG systems are not only accurate but also complete and directly address the user’s query without including unnecessary or redundant information

# EVALUATING RAG SYSTEM: FAITHFULNESS (GROUNDEDNESS) (1/2):ANSWER/CONTEXT

Generation component: synthesizes responses based on the retrieved data

Faithfulness  
(Groundedness)

- Assess the **factual alignment** and semantic similarity between the model's **response and the retrieved documents**, ensuring that the generated answer is **contextually appropriate** and **factually grounded** in the retrieved information.
  - How factually accurate is the generated answer?
  - Can all the claims made in the model's answer be directly inferred from the given context?
- Evaluation process: hybrid – automated system (CoT prompting; semantic matching) + human judgement
  - 1. Identify **a set of claims** made in the generated answer
  - 2. Cross-check each claim against the given context to determine its factual consistency and whether it can be **directly inferred from the retrieved information**
  - 3. Calculate the Groundedness score using the following formula (0 – 1)

$$\text{Groundedness} = \frac{\text{Number of claims inferable from the given context}}{\text{Total number of claims in the generated answer}}$$

## EVALUATING RAG SYSTEM: FAITHFULNESS (GROUNDEDNESS) (2/2) :ANSWER/CONTEXT

### Example

- Question: **“Where and when was Einstein born?”**
- Context: **“Albert Einstein (born March 14, 1879) was a German-born theoretical physicist.”**
- Answer:
  - **High groundedness:**“Einstein was born in **Germany** on **March 14<sup>th</sup>, 1879.**” (aligns with the factual information provided in the context)
  - **Low groundedness:**“Einstein was born in Germany on **March 20<sup>th</sup>, 1879.**” (contradicts the date of birth mentioned in the context)

The Groundedness metric is crucial in ensuring the reliability and trustworthiness of responses generated by RAG systems, as it directly relates to the accuracy and factual consistency of the information provided in response to a user's query.

## EVALUATING RAG SYSTEM: CONTEXT RELEVANCE (1/2): QUERY/CONTEXT

**Retrieval** component: responsible for fetching **relevant information** from a (vector) database or corpus in response to a user's query

### Context Relevance

- Measures the alignment of the **retrieved information** with the user's query, ensuring that only **essential information** is included to address the query effectively.
  - Is the **retrieved context pertinent and appropriate** for effectively addressing the given query?
  - Is **only essential and directly relevant** information included?
- Evaluation process: a two-step procedure
  - 1. Sentence-level Relevance Scoring: Each sentence in the retrieved context is assigned a relevance score based on its semantic similarity to the query, using measures like cosine similarity between embeddings.
  - 2. Overall Context Relevance Quantification: The final context relevance score is calculated as the ratio of relevant sentences to the total number of sentences in the retrieved context (0 – 1)

$$\text{Context Relevance} = \frac{\text{Number of relevant sentences within the retrieved context}}{\text{Total number of sentences in retrieved context}}$$

## EVALUATING RAG SYSTEM: CONTEXT RELEVANCE(2/2): QUERY/CONTEXT

### Example

- Question: “**What is the capital of France?**”
- Context: “*France, in Western Europe, encompasses medieval cities, alpine villages and Mediterranean beaches. **Paris, its capital**, is famed for its fashion houses, classical art museums including the Louvre and monuments like the Eiffel Tower.*”
  - **High context relevance:**“*France, in Western Europe, encompasses medieval cities, alpine villages and Mediterranean beaches. **Paris, its capital**, is famed for its fashion houses, classical art museums including the Louvre and monuments like the Eiffel Tower.*” (includes information **directly mentioning** Paris as the capital)
  - **Low context relevance:**“*The country is also **renowned for its wines and sophisticated cuisine**. Lascaux's **ancient cave drawings**, **Lyon's Roman theater** and the **vast Palace of Versailles** attest to its rich history.*” (includes additional, tangential information)

By evaluating Context Relevance, the RAG system can ensure that the retrieved information is concise, focused, and directly addresses the user's query, enhancing the efficiency and accuracy of the generated response.

---

## LLM PREFERENCE MEASUREMENT



Insights from Measurement Theory



Offline Metric: Measurement Scale Options



Pros and Cons: MaxDiff as Next Generation



MaxDiff Example: Answer Card Optimization

---

# MEASUREMENT SCALES

## Nominal (Categorical)

- Classifies data into **distinct groups without any inherent order**.
- Example: Gender, Eye Color.
- Statistical Operations: Counting, **mode** calculation.

## Ordinal

- Data is **ranked in a meaningful order**, but **intervals between values are not uniform**.
- Example: Education Levels, Customer Satisfaction, Rank-order
- Statistical Operations: **Median**, mode, **non-parametric tests**.

## Interval

- Numeric scale with **equal intervals** but **no true zero point**.
- Example: Temperature in Celsius/Fahrenheit.
- Statistical Operations: **Mean**, median, mode, **addition/subtraction**, parametric tests.

## Raio

- Numeric scale with **equal intervals** and an **absolute zero**.
- Example: Weight, Height, Age, Income.
- Operations: All arithmetic operations, including **multiplication and division**.

## OFFLINE METRIC - LLM MEASUREMENT SCALE OPTIONS

### Standard Rating

- Also called as Likert Scales
- Widely Used Before / Less Used Now
- Simple but suffers from bias and lack of differentiation.

Example:

Rate each feature from 1-5.

### Pairwise Comparison

- Also called as “comparative eval”
- Industry Standard (From 2022) (OpenAI, Anthropic)
- Precise but impractical for large sets.

Example:

Do you prefer A or B?

### MaxDiff

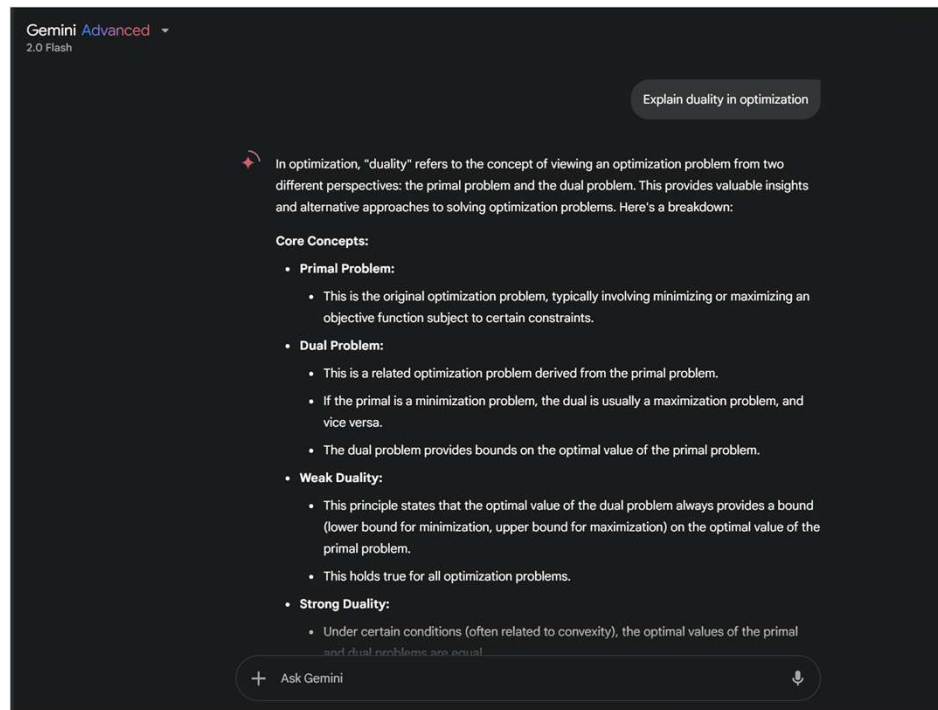
- Maximum Difference Scale
- Also called as Best Worst Scaling
- Reduces bias, handles large lists efficiently, and delivers clear, actionable insights.
- **Not explored yet** for LLM Eval Application

Example:

Among A, B, C, D – pick the most preferred & least preferred.



# STANDARD RATING (5-POINT LIKERT SCALE) - STATED



How satisfied was the user?



Very Satisfied



Satisfied



Neutral



Dissatisfied



Very dissatisfied

## MORE DISCRIMINATION FROM TRADE-OFF QUESTIONS



User Preference

---

## EXAMPLE: NEW JOB CHOICE

**Location**

**Salary**

**Option 1**



**\$70K**

**Option 2**



**\$100K**

User Preference

---

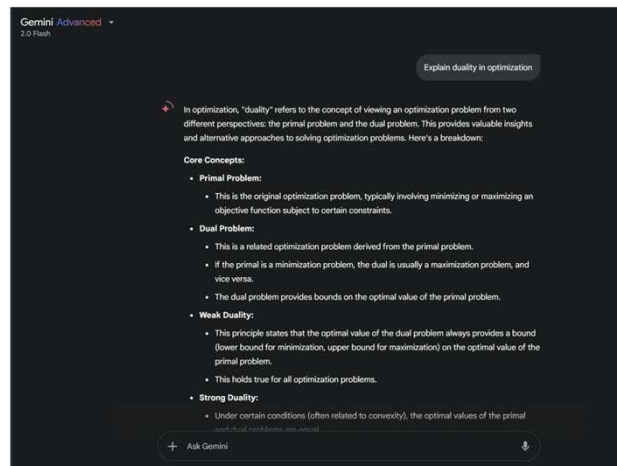
## COMPARATIVE EVAL REVEALS WHAT PEOPLE REALLY WANTS

<b>Job Attribute</b>	<b>What MBA's Said... (Stated Rating)</b>	<b>Rank Based on Comparative Eval</b>
Salary	6th	1st
Region of US	2nd	2nd
Job Location	5th	3rd
People/Culture	1st	4th
Functional Area	3rd	5th
Firm Growth	7th	6th
Business Travel	8th	7th
Opp. to Advance	4th	8th

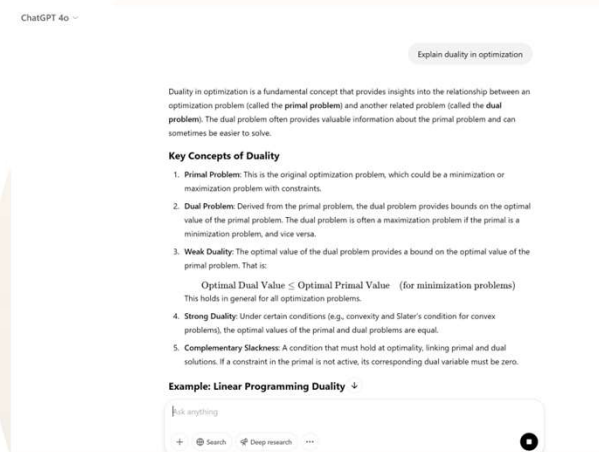
User Preference

# PAIR-WISE COMPARISON: GEMINI VS. CHATGPT 4O

## Version A: Gemini 2.5



## Version B: GPT-4o



Which chatbot response do you prefer?



A >> B



A > B



A == B



A < B



A << B

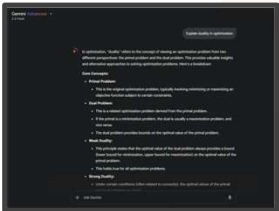
# MAXDIFF: GEMINI, CHATGPT 4O, GROK 3

Based on Chatbot responses, please choose your Favorite and Least Favorite

Task 1/7

Favorite

Least Favorite



---

## PROS AND CONS

### Pros

### Cons

#### Standard Rating

- Simple and intuitive for respondents.

- Limited discrimination between items; susceptible to biases; hard to answer

#### Pairwise Comparison

- Intuitive, Provides precise preference data.

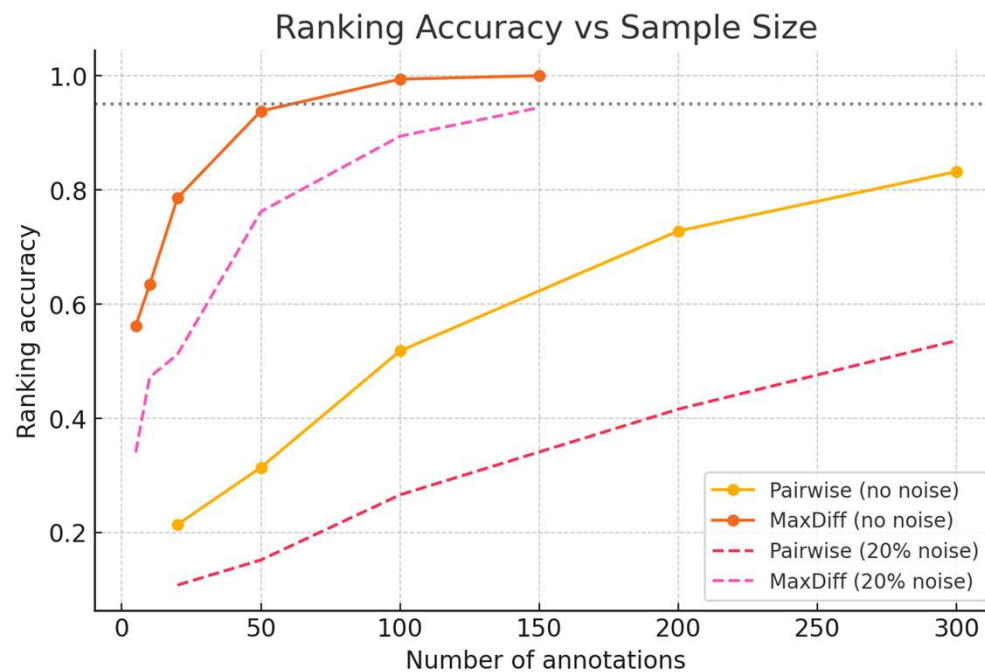
- Impractical with large item sets; can lead to respondent fatigue.

#### MaxDiff

- Enhances discrimination by forcing trade-offs; reduces biases; allows more robust optimization beyond measurement

- More complex to design and analyze; may be cognitively demanding.

# MAXDIFF IS MORE SAMPLE EFFICIENT AND ROBUST: SIMULATION RESULTS



5 Model Candidates: showing 4 candidates per each task for MaxDiff



---

## MAXDIFF EXAMPLE:ANSWER CARD OPTIMIZATION

(1) Key Attributes to Consider:

- **Presentation Format**:Text, Rich details with picture,Text with box.
- **Font Size**: Small, Medium, Large.
- **Color Scheme**: Light theme, Dark theme, Colorful theme.
- Implementing MaxDiff for Answer Card Optimization:

(2) Define the combinations: options to show

(3) Design the MaxDiff survey: Sawtooth software

(4) Conduct MaxDiff survey question: Sawtooth software

(5) Analyzing the results: Hierarchical Bayes to Retrieve underlying preference

---

## CHATBOT ARENA / MAXDIFF RESOURCE

- Chatbot Arena (LMSYS): Crowd-based pair-wise comparison
  - <https://lmarena.ai/>
- MaxDiff Exercise YouTube Video: Sawtooth Software
  - <https://youtu.be/Sg-cxgCOVf4>
- MaxDiff Explainer from Sawtooth Software
  - <https://sawtoothsoftware.com/maxdiff>
- Interactive Apps and Data Collection and Voting Results
  - Gradio: <https://www.deeplearning.ai/short-courses/building-generative-ai-applications-with-gradio/>
  - Chatbot Response Collector: Claude.ai: [file:///C:/Users/h\\_min/Downloads/prompt-collector-MVP.html](file:///C:/Users/h_min/Downloads/prompt-collector-MVP.html)