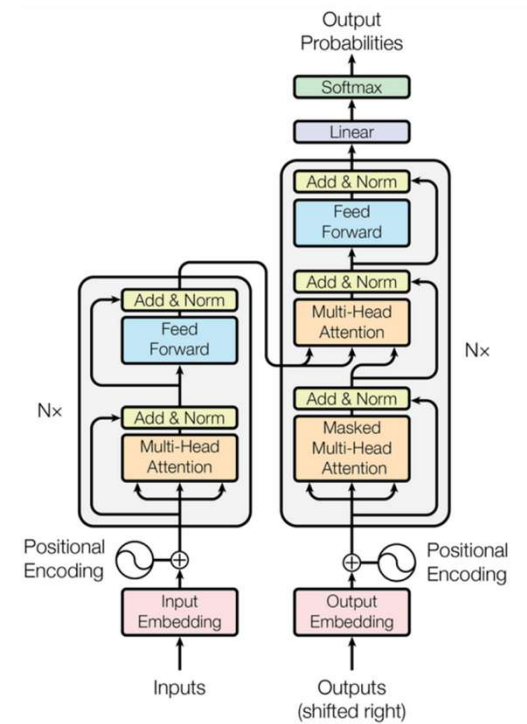


LARGE LANGUAGE MODEL AS REINFORCEMENT LEARNING AGENT

Minha Hwang



TWO STAGES OF LARGE LANGUAGE MODEL DEVELOPMENT

Pre-Training: Base LLM (GPT3)

predict next word, based on text training data

- **Self-supervised**

Once upon a time, there was a unicorn
that lived in a magical forest with
all her unicorn friends

What is the capital of France?
What is France's largest city?
What is France's population?
What is the currency of France?

Post-Training: Instruction Tuned LLM (ChatGPT)

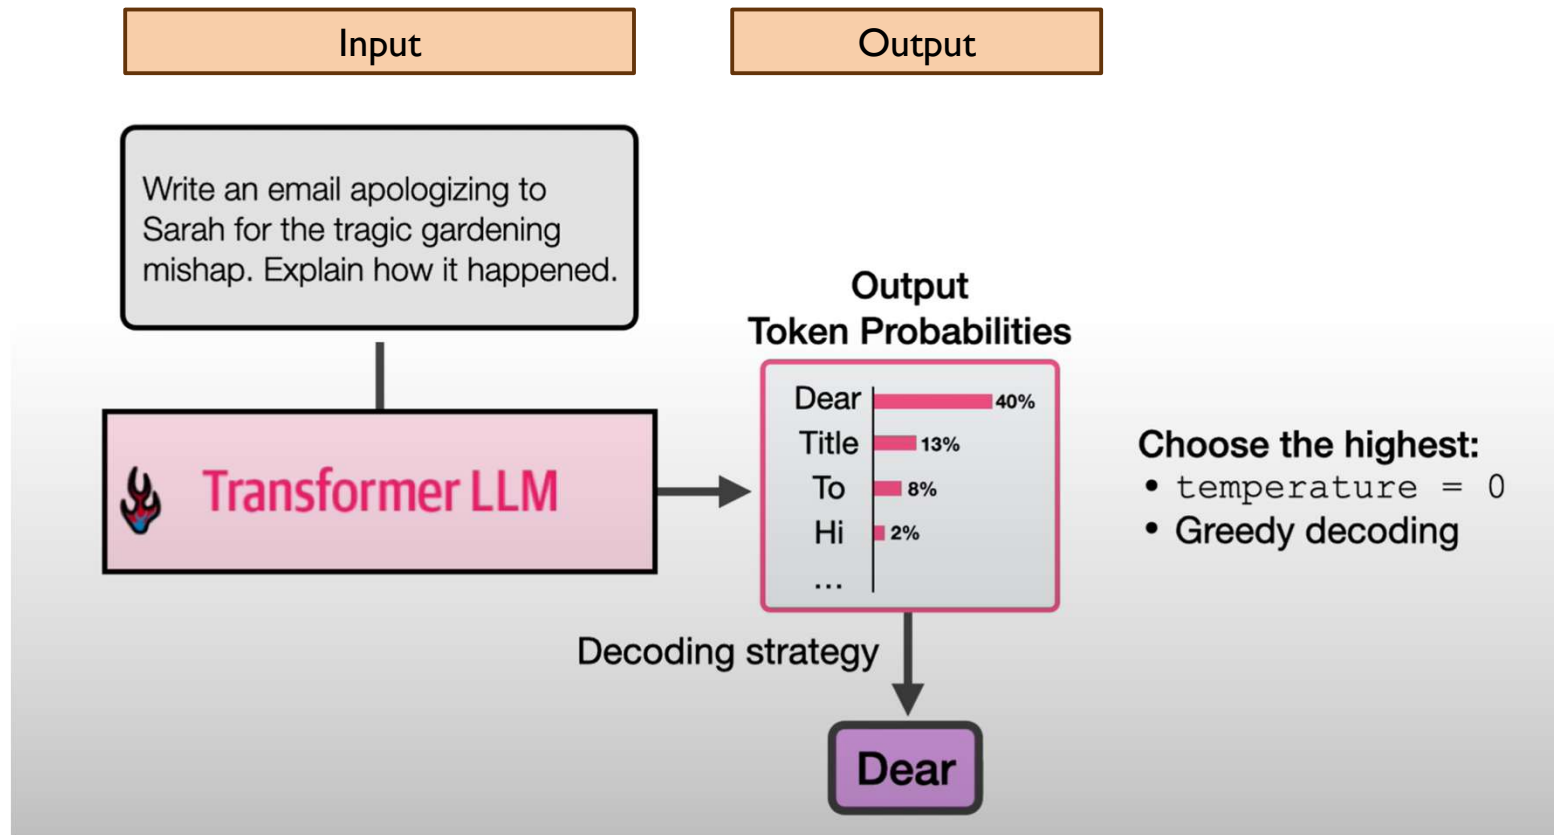
Tries to **follow instructions**

Fine-tune on **instructions and good response pairs**

- **Human labeled data**: Instruction – Response Pair
- **SFT**: Supervised Fine-tuning
- **RLHF** (Reinforcement Learning with Human Feedback) or **DPO** (Direct Preference Optimization)

What is the capital of France?
The capital of France is Paris.

TRANSFORMER LLM: INPUT AND OUTPUT



AUTOREGRESSIVE LANGUAGE MODELS: PROBABILITY KERNEL

- A language model is a **probability kernel μ** given a prefix of words: $\mu: X \rightarrow Pr(Y)$
 - Stochastic in nature: **A same prefix X** can give a **random output** sampled from a probability distribution μ_X (i.e., generative) → A key reason for factual inaccuracy, inconsistency or hallucination (making stuff up)
- A language model calculates **$Pr(s)$** given a sequence of words: **$s = (w_1, w_2, \dots, w_{T-1}, w_T)$**
- An autoregressive language model calculates this **conditional on a previous sequence of words**:

$$\begin{aligned} Pr(s) &= Pr(w_1, w_2, \dots, w_{T-1}, w_T) \\ &= \prod_{t=1}^T Pr(w_t | w_1, w_2, \dots, w_{t-1}) \end{aligned}$$

- **Next-word prediction**: Given a prefix $(w_1, w_2, \dots, w_{t-1})$, calculate the probability of the next word w_t (Conceptually same to time series with path dependence)

Source: Prof. Kyunghyun Cho

AUTOREGRESSIVE LANGUAGE MODELS: SIMPLE EXAMPLE

- 4-word sentence example: “I am a student”

$$Pr(s) = Pr(w_1, w_2, w_3, w_4) = Pr(w_1) \times Pr(w_2|w_1) \times Pr(w_3|w_1, w_2) \times Pr(w_4|w_1, w_2, w_3)$$

- All you need is “counting” (if there are large amounts of data)

$$Pr(w_2|w_1) = \frac{count(w_1, w_2)}{count(w_1)} \longrightarrow \text{2-grams (Bigrams)}$$

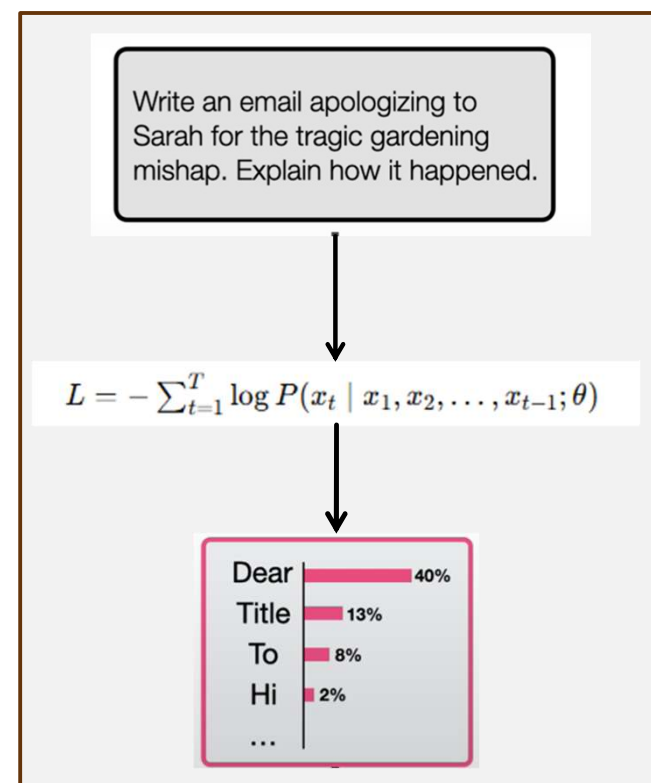
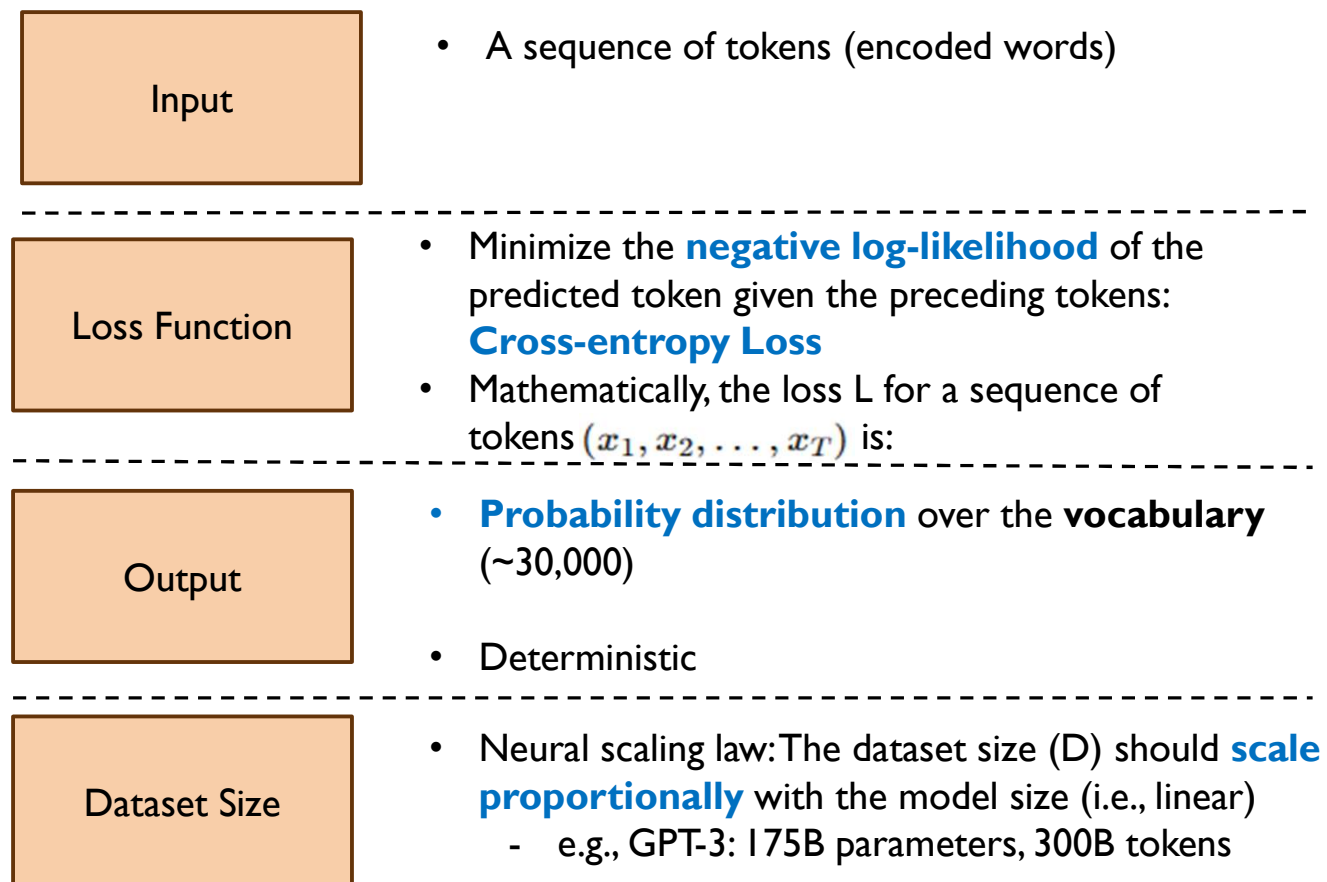
$$Pr(w_3|w_1, w_2) = \frac{count(w_1, w_2, w_3)}{count(w_1, w_2)} \longrightarrow \text{3-grams (Trigrams)}$$

$$Pr(w_4|w_1, w_2, w_3) = \frac{count(w_1, w_2, w_3, w_4)}{count(w_1, w_2, w_3)} \longrightarrow \text{4-grams}$$

- Problems:
 - This requires **a lot of space (RAM)**
 - Count-based language models **cannot generalize**: A certain sentence **does not appear** in the corpus

Source: Prof. Kyunghyun Cho

PRE-TRAINING: SELF-SUPERVISED LEARNING



PRE-TRAINED LLM EVAL - PERPLEXITY: PREDICTIVE ACCURACY (2/2)

Example Calculation: “The cat sat on the mat”

- $P(\text{“The”}) = 0.2$
- $P(\text{“cat”} | \text{“The”}) = 0.1$
- $P(\text{“sat”} | \text{“The cat”}) = 0.15$
- $P(\text{“on”} | \text{“The cat sat”}) = 0.3$
- $P(\text{“the”} | \text{“The cat sat on”}) = 0.25$
- $P(\text{“mat”} | \text{“The cat sat on the”}) = 0.05$

First, calculate the average negative log probability:

$$-\frac{1}{6} (\log(0.2) + \log(0.1) + \log(0.15) + \log(0.3) + \log(0.25) + \log(0.05)) \approx 1.8992$$

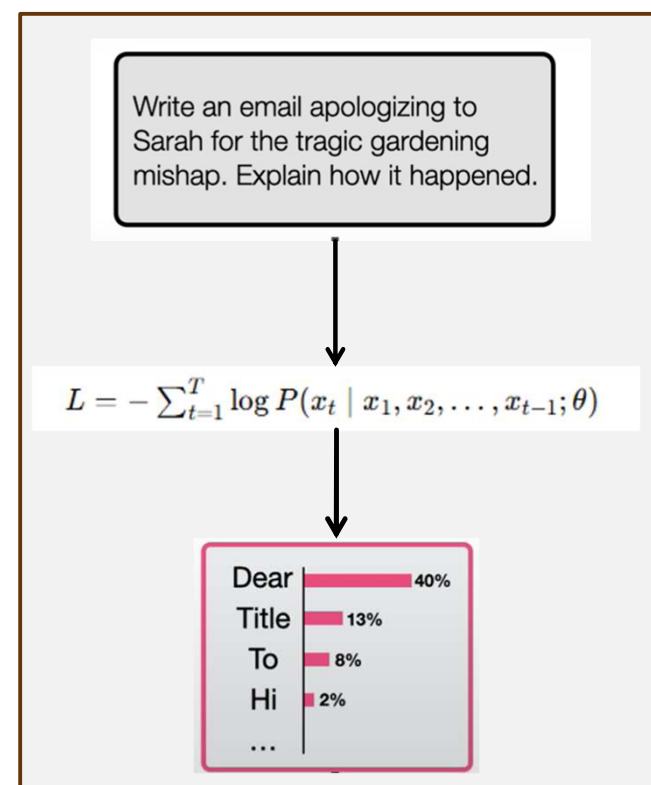
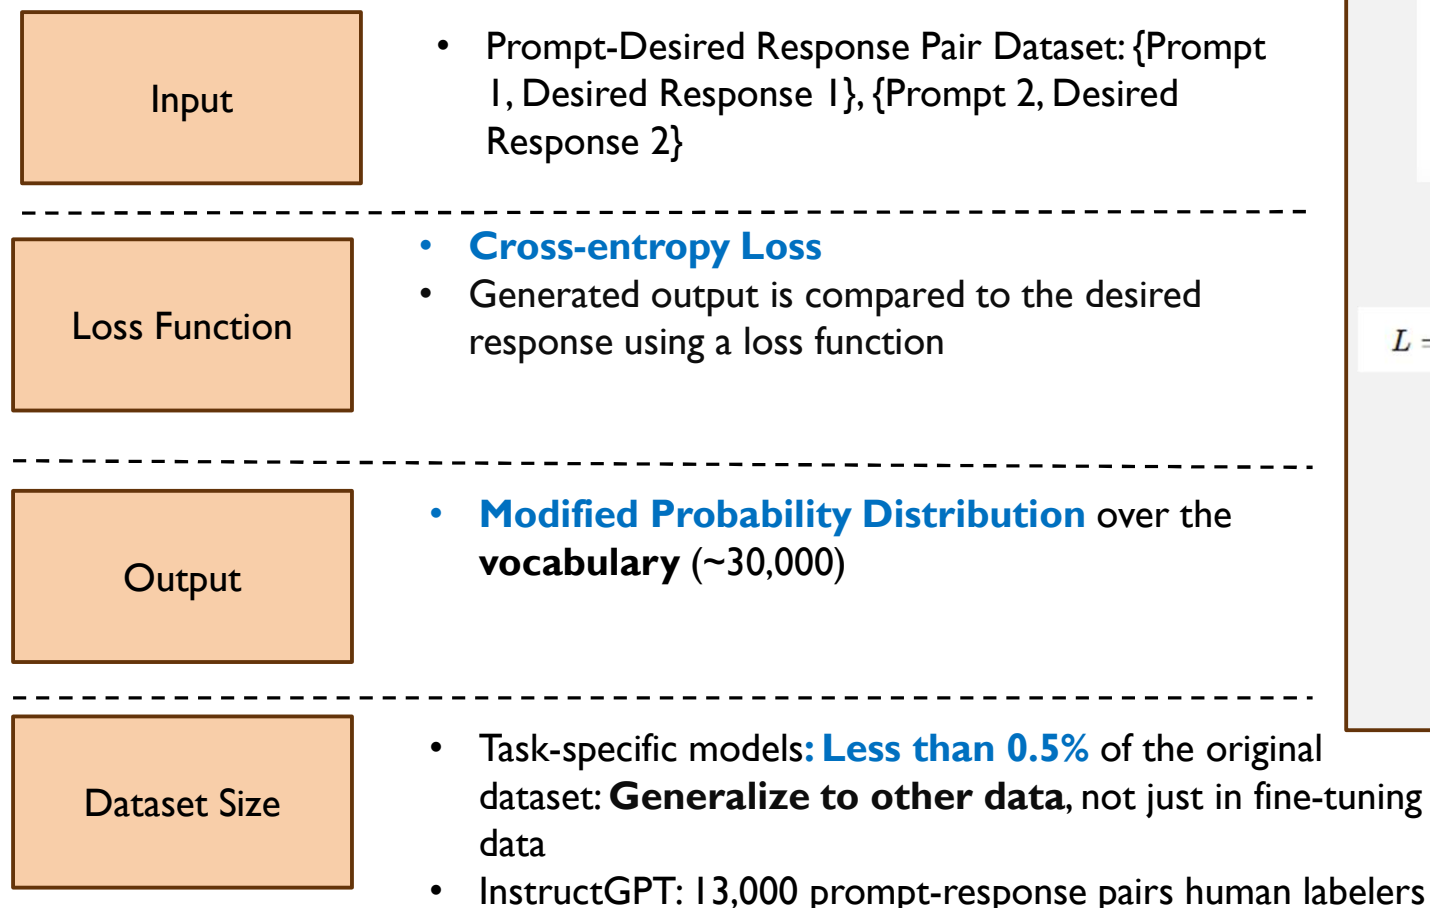
Then, exponentiate to find perplexity:

$$PP = \exp(1.8992) \approx 6.68$$


This means the model, on average, considers about 6.68 possible next words, indicating its uncertainty in prediction.

$$\text{Perplexity}(PP) = \exp \left(-\frac{1}{N} \sum_{i=1}^N \log P(w_i | w_{<i}) \right)$$

POST-TRAINING: SUPERVISED FINE TUNING



REINFORCEMENT LEARNING: AGENT AND ENVIRONMENT




- At each step t , **the agent**:
 - Execute action $A(t)$
 - Receives observation $O(t)$
 - Receives scalar reward $R(t)$
- The **environment**:
 - Receives action $A(t)$
 - Emits observation $O(t+1)$
 - Emits scalar reward $R(t+1)$
- **t increment** at environment step

- **Sequential Decision Making**

- Reward Hypothesis: All goals can be described by the maximization of **expected cumulative reward (scalar)**

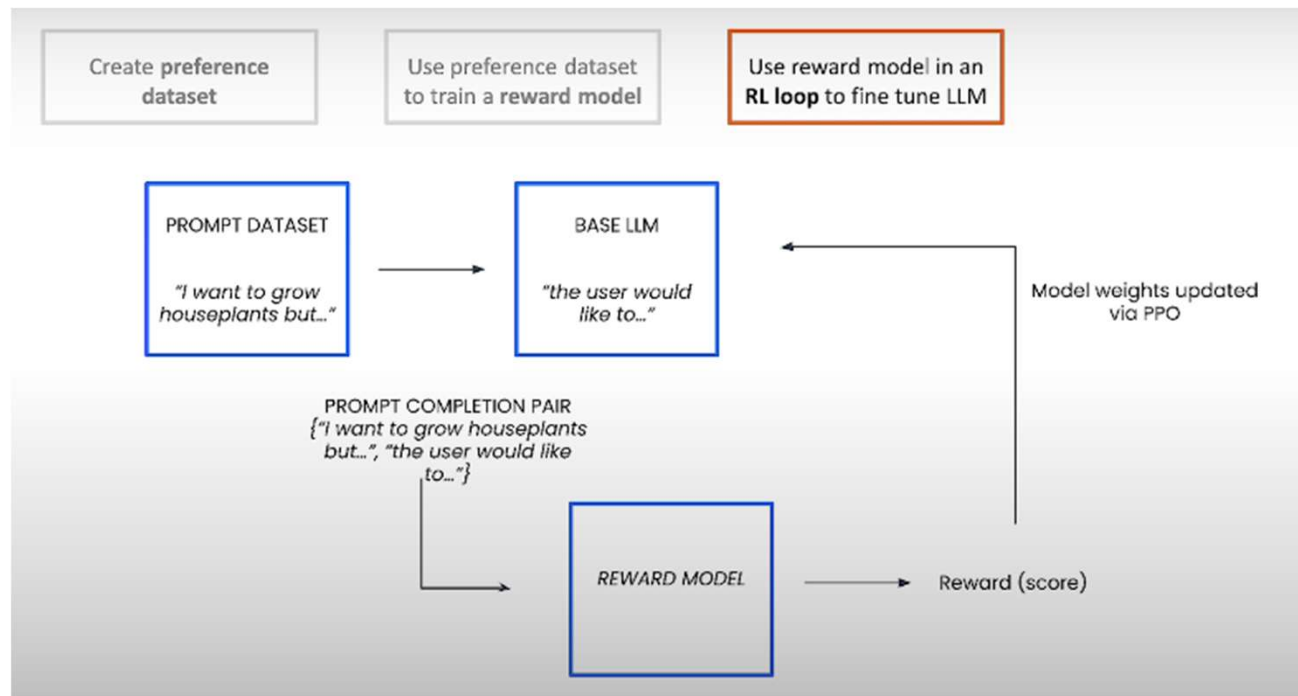
POST-TRAINING: RLHF TO FINE TUNE LLM



- Proximal Policy Optimization (PPO)
- DeepSeek: GRPO

Source: <https://learn.deeplearning.ai/courses/reinforcement-learning-from-human-feedback>

POST-TRAINING: RLHF



- **Sequential Decision Making**
- Reward Hypothesis: All goals can be described by the maximization of **expected cumulative reward (scalar)**

POST-TRAINING: RLHF – REWARD MODEL – TRAINING (1/3)

Input

- **Preference Dataset** - Pairwise {Prompt, Winning Candidate, Losing Candidate, Choice}
- Annotated by Human (Subjective)

Loss Function

- Minimize **Pairwise Loss**

$$\mathcal{L}(\theta) = -\frac{1}{\binom{K}{2}} \sum_{(x, y_w, y_l)} \log(\sigma(r_\theta(x, y_w) - r_\theta(x, y_l)))$$

Here:

- x is the prompt.
 - y_w and y_l are the preferred and less preferred responses, respectively.
 - $r_\theta(x, y)$ is the reward model's score for a given prompt-response pair.
 - σ denotes the sigmoid function.
 - K is the number of responses ranked by human annotators for each prompt.
- 10K – 100K range
 - InstructGPT:
 - Reward Model Dataset: ~ 33,000 examples. Human labelers ranked multiple responses to the same prompt

Dataset Size

POST-TRAINING: RLHF – REWARD MODEL – INFERENCE (2/3)

Input


- **Prompt Dataset** - {Prompt, Response}

Model

- Reward Model (LLM), trained with **Preference Dataset**: {Prompt, Winning Candidate, Losing Candidate, Choice}

Output

- Scalar indicating how good response is



POST-TRAINING: RLHF – PROXIMAL POLICY OPTIMIZATION (3/3)

Input

- **Proximal Policy Optimization (PPO) Dataset:** InstructGPT ~31,000 prompts used to generate responses without human intervention during training, {Prompt}

Loss Function

Policy Optimization: The language model is fine-tuned using reinforcement learning to maximize the rewards predicted by the reward model. A common approach is to use Proximal Policy Optimization (PPO) with a loss function that balances achieving high reward and maintaining the model's output distribution close to the original model to prevent divergence:

$$L(\phi) = \mathbb{E}_{(x,y) \sim D_{\pi_\phi}} \left[r_\theta(x,y) - \beta \log \left(\frac{\pi_\phi(y|x)}{\pi_{\text{SFT}}(y|x)} \right) \right]$$

In this equation:

- π_ϕ is the policy of the fine-tuned model.
- π_{SFT} is the policy of the supervised fine-tuned model before reinforcement learning.
- β is a scaling factor that controls the strength of the penalty for deviating from the original policy.

CONNECTING RL AND LLMS (1/2)

Connection in LLM

Sequential
Decision Making

- Next token prediction to maximize prediction accuracy (pre-training) and rewards (post-training) over vocabulary (discrete set)

State

- Input prompt and previous generated tokens

Action

- Token chosen from the vocabulary

Reward

- Prediction accuracy + rewards from aligning with preference (weighted)

$$r(s_t, y_t) = \lambda_1 \log \pi(y_t | s_t) + \lambda_2 \text{Metric}(s_t, y_t),$$

Source: [Large Language Models as Reinforcement Learning Agents in Token Space: A Theoretical Framework by Miquel Noguer I Alonso :: SSRN](#)

CONNECTING RL AND LLMS (2/2): APPLICABILITY OF RL RESEARCH

Description

Hierarchical Decision

- Motivated by AlphaStar's multi-scale decision-making
- High-level planning and detailed token generation

Self-dialogue

- Motivated by self-play training in AlphaGo, AlphaStar
- Self-dialogue: Models can engage in dialogue, critiquing and improving each other's outputs

Adaptive decoding

- Develop decoders that balance exploration and exploitation based on state uncertainty

Hybrid Models

- Combine maximum likelihood training with RL-based fine-tuning

Source: [Large Language Models as Reinforcement Learning Agents in Token Space: A Theoretical Framework by Miquel Noguer I Alonso :: SSRN](#)

Reference

- **Transformer Large Language Model**

- <https://learn.deeplearning.ai/courses/how-transformer-llms-work>
- <https://learn.deeplearning.ai/courses/attention-in-transformers-concepts-and-code-in-pytorch>

- **Neural Sampling Law**

- [\[2203.15556\] Training Compute-Optimal Large Language Models](#)

- **Reinforcement Learning with Human Feedback**

- [InstructGPT: \[2203.02155\] Training language models to follow instructions with human feedback](#)
- [Paper page - Fine-Tuning Language Models from Human Preferences](#)

- **Large Language Model and Reinforcement Learning Connection**

- [Large Language Models as Reinforcement Learning Agents in Token Space: A Theoretical Framework by Miquel Noguer I Alonso :: SSRN](#)
- [\[2306.07929\] Large Language Models Are Semi-Parametric Reinforcement Learning Agents](#)
- [WindyLab/LLM-RL-Papers: Monitoring recent cross-research on LLM & RL on arXiv for control. If there are good papers, PRs are welcome.](#)
- [\[2106.01345\] Decision Transformer: Reinforcement Learning via Sequence Modeling](#)
- [kzl/decision-transformer: Official codebase for Decision Transformer: Reinforcement Learning via Sequence Modeling.](#)